

Setting up Google Cloud and BigQuery

This is the step-by-step guide to set up Google Cloud and Google BigQuery as part of the Stroopwafelshop project. It contains the following steps:

1. Creating a Google Cloud account
2. Creating a project in Google Cloud
3. Creating a BigQuery dataset
4. Creating a service account for Airbyte with permission to access BigQuery
5. Creating additional service accounts for dbt Cloud and Tableau

Let's get to work!

Creating a Google Cloud account

To use BigQuery you will need to set up an **account** in Google Cloud. Google Cloud charges for the use of many of its services, including BigQuery. However, BigQuery offers a generous monthly **free tier** which should be more than enough to follow along in the chapter (read more on <https://cloud.google.com/free/docs/free-cloud-features#bigquery>). If you already have an account in Google Cloud, you should be able to make use of this offering and do not need to create a new account.

If you do not own an existing Google Cloud account, you can create one and get \$300 dollars in free **credits** during a 90-day trial period, after which there are no costs unless the user explicitly upgrades from the **trial mode** to a **full account** (as of February 2024, shown in **Figure 1**). To create this trial account, you will need to provide credit card details, although no money will be charged, and you do not need to spend money during or after following this chapter.

If you already have an existing Google Cloud account with valid payment information, you can use that and skip the account creation steps described below. However, if you have an account but do not have usable credits, you can still create a new account.

Unfortunately, there is no way to make use of the full features of Google BigQuery which we need for the chapter, without providing payment information. As stated, you will not be charged anything as long as you don't manually upgrade the account, which is not necessary for following the chapter. Furthermore, Google Cloud is the only tool that we will use that requires you to give any payment details at all.

Note

Google also offers a totally free version of BigQuery called **BigQuery Sandbox**. You can use it without providing payment information. However, this sandbox environment is more limited and does not support certain features that are required by Airbyte. This is why we still have to provide payment information to access the full capabilities of BigQuery. Read more on BigQuery Sandbox here: <https://cloud.google.com/bigquery/docs/sandbox>.

If you want to create a new account: visit the page on <https://cloud.google.com/free> and fill in your details. After filling in the details, Google verifies the payment details and might require confirmation. Next, you should end up in the home screen where you are greeted with some information on your account (as shown in **Figure 3**).

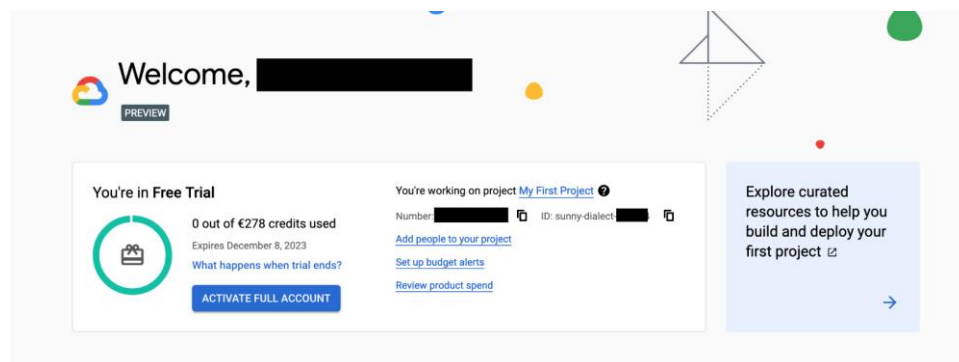


Figure 1 - Google Cloud welcome screen

As you can see, you are using the **Free Trial** mode, with the full amount of credits remaining. The big blue button **Activate Full Account** indicates the manual step necessary to upgrade your account. *As long as you do not click (and confirm) that, you will not be charged anything.* So, finger off the trigger button! Furthermore, you can read on the page that you are working on a project called **My First Project**, which is the default **project** created for you.

Creating a Project

A project is the fundamental building block in Google Cloud, and most services need to be associated with a project before you can use them. **billing**, **access management**, and **resource** creation are all tied to a project, and it is possible to create multiple projects in an account and switch between them.

Create a new project called **stroopwafelshop**. Next, we will set up BigQuery for data loading.

Setting up BigQuery

Navigate to <https://console.cloud.google.com/bigquery>. You should be greeted by BigQuery's welcome page, as shown in **Figure 2**.

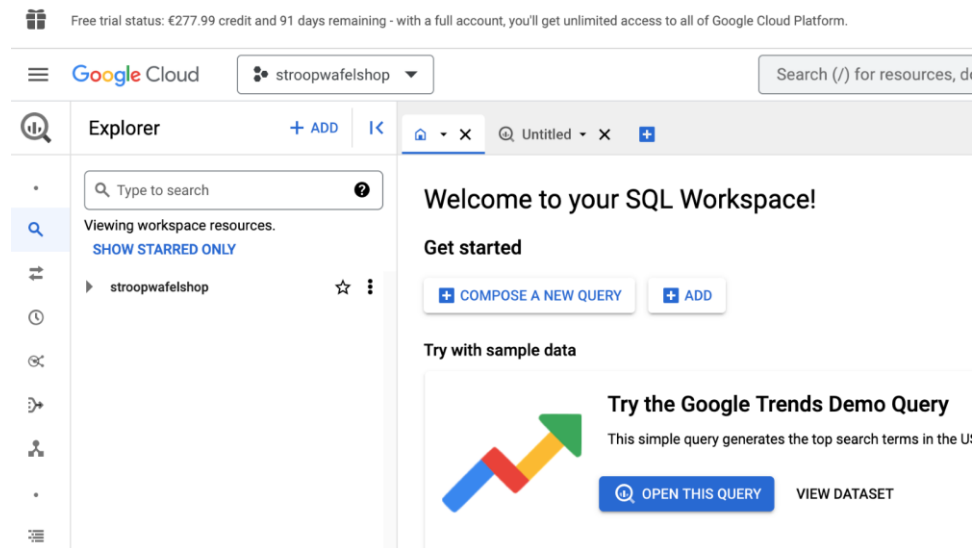


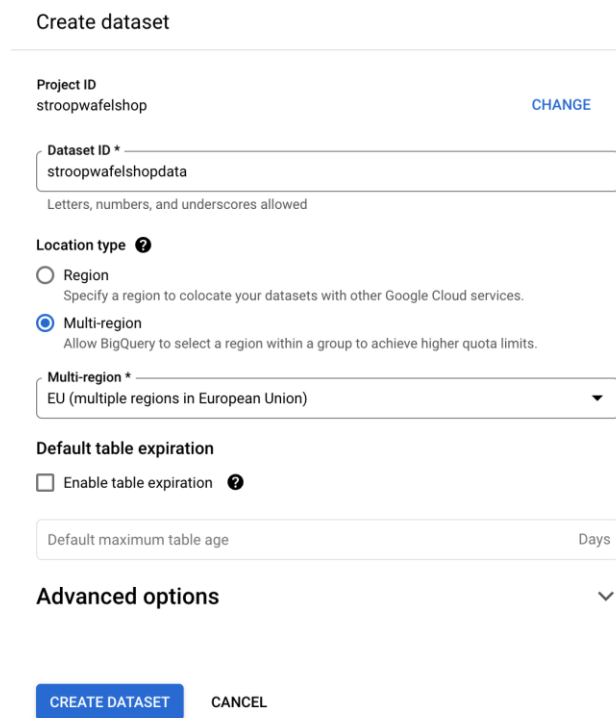
Figure 2 - Google BigQuery landing page

This shows the **SQL Workspace**, where you can import data to create tables and then use SQL to query them. BigQuery also has many other features, such as exploring public datasets, transferring data from various sources, scheduling queries, etc. We will stick to the basics by creating tables from raw data and modeling them using SQL. First, we will need to create a **dataset**.

Creating a BigQuery dataset

In BigQuery, a dataset is a container for **tables**, **views**, and other data structures. It provides a way to organize and control access to these structures, much like how a folder organizes files on your computer. We want to create a dataset to store the source data.

Choose the option to create a new dataset, and name it **stroopwafelshopdata**. When creating a dataset you must specify a **location** for the storage. This can be a specific region such as **europa-west1**, or a **multi-region** location such as the **US**. Since the Stroopwafelshop is in the EU, we will select **EU**, as shown in **Figure 3**. Leave the other options as they were and select **CREATE DATASET**.



The screenshot shows the 'Create dataset' form in the Google Cloud BigQuery console. The form is titled 'Create dataset' and has a 'Project ID' field with the value 'stroopwafelshop' and a 'CHANGE' link. Below this is a 'Dataset ID *' field with the value 'stroopwafelshopdata' and a note 'Letters, numbers, and underscores allowed'. The 'Location type' section has two options: 'Region' (unselected) and 'Multi-region' (selected). The 'Multi-region' option has a note 'Allow BigQuery to select a region within a group to achieve higher quota limits.' Below this is a 'Multi-region *' dropdown menu with the value 'EU (multiple regions in European Union)'. The 'Default table expiration' section has a checkbox 'Enable table expiration' (unchecked) and a note. Below this is a 'Default maximum table age' field with the value 'Days'. The 'Advanced options' section is collapsed. At the bottom are two buttons: 'CREATE DATASET' and 'CANCEL'.

Figure 3 – Creating the dataset

After creating the dataset inside of your project, it should be visible on the left-hand side of the screen. Next up, we will create **service accounts** that are allowed to access the dataset.

Creating service accounts

When using Google Cloud, you can navigate to the webpage using the browser as we've been doing, and interact with other services from there. This part of Google Cloud is called the **console** (<https://cloud.google.com/cloud-console>).

Before being able to use the console, you must authenticate by logging in. Once logged in, Google takes care of all authentication and authorization. But if you want to connect to these services using a **third-party** application or software, you will need to authenticate with Google's services.

Essentially, whenever an application such as Airbyte wants to communicate with Google's services over the internet, they will first have to prove that they are allowed access to your project in Google Cloud and be able to read your data. Generally, this is done by creating special non-user accounts. These accounts are commonly known **service accounts**. A service account can be assigned granular permissions. For instance, they could be permitted to only use a single service, such as BigQuery. This can help reduce the potential damage in case of a compromised account. Therefore, it is best practice to create a separate service account for each application that needs to access the cloud platform, and only assign them permissions they need, not more. Considering this, let's create the service account.

In Google Cloud, visit the **IAM (Identity and Access Management)** page. This is where the permissions in your project can be managed. You can either search for the service in the search bar or access it from <https://console.cloud.google.com/iam-admin/iam>. In the sidebar on the left, choose **Service Accounts**. This will show all existing service accounts (as shown in **Figure 4**). Currently you do not have any.

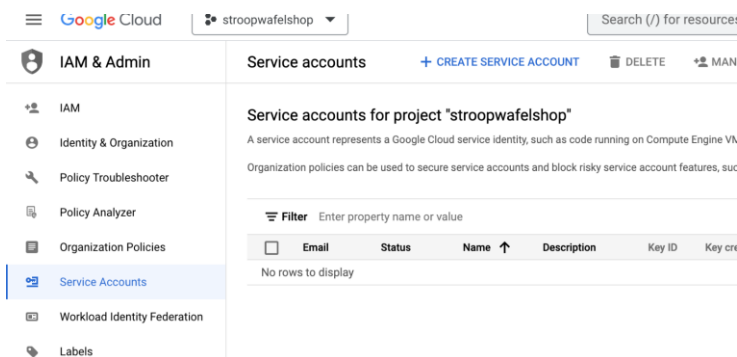


Figure 4 - Service accounts overview

Create a service account called **airbyte**. Once created, a unique identifier in the form of an email address is generated in the form of name@project.iam.gserviceaccount.com. In our

case, it will be airbyte@stroopwafelshop.iam.gserviceaccount.com. Add a description to specify its use, as shown in **Figure 5**.

The screenshot shows the 'Service account details' form in the Google Cloud console. It includes fields for 'Service account name' (filled with 'airbyte'), 'Service account ID' (filled with 'airbyte'), and 'Service account description' (filled with 'Used by Airbyte to connect to BigQuery'). The email address is automatically generated as 'airbyte@stroopwafelshop.iam.gserviceaccount.com'. A 'CREATE AND CONTINUE' button is at the bottom of the first section. Below this, there are two optional sections: 'Grant this service account access to project (optional)' and 'Grant users access to this service account (optional)'. At the very bottom are 'DONE' and 'CANCEL' buttons.

Figure 5- Creating the service account

Next, select the option to grant access to a project. We will grant permission to the service account to access our project, and specifically, BigQuery. The easiest way for this is to select a pre-defined **role**, as shown in **Figure 6**.

Roles in Google Cloud act as a bundle of pre-defined permissions, defined by Google for a specific workload. In our case, we need the **BigQuery User** and **BigQuery Data Editor** roles to interact with the BigQuery service from Airbyte. For more information on this role, see <https://cloud.google.com/bigquery/docs/access-control> and <https://docs.airbyte.com/integrations/sources/bigquery#service-account>). This will allow read and write access to the BigQuery datasets in the project.

Additionally, we could allow access to specific datasets in the BigQuery console, but for now, this is not required. Select both roles, and finally create the account.

✓ Service account details

2 Grant this service account access to project (optional)

Grant this service account access to stroopwafelshop so that it has permission to complete specific actions on the resources in your project. [Learn more](#)

Role
BigQuery User

IAM condition (optional) ?
+ ADD IAM CONDITION

When applied to a project, access to run queries, create datasets, read dataset metadata, and list tables.
When applied to a dataset, access to read dataset metadata and list tables within the dataset.

Role
BigQuery Data Editor

IAM condition (optional) ?
+ ADD IAM CONDITION

Access to edit all the contents of datasets

+ ADD ANOTHER ROLE

CONTINUE

3 Grant users access to this service account (optional)

DONE CANCEL

Figure 6 - Selecting the BigQuery User and BigQuery Data Editor roles

This should bring you back to the **Service Accounts** page where you started, but with the newly created service accounts' email address visible in the middle of the screen as shown in **Figure 10**. Next, we need to create a key to serve as **credentials** for our service account.

Service accounts for project: stroopwafelshop

A service account represents a Google Cloud service identity, such as code running on Compute Engine VMs, App Engine apps, or systems running outside Google. [Learn more about service accounts](#)

Organization policies can be used to secure service accounts and block risky service account features, such as automatic IAM Grants, key creation/upload, or the creation of service accounts entirely. [Learn more about service account organization policies](#)

Filter	Email	Status	Name	Description	Key ID	Key creation date	OAuth 2 Client ID	Actions
	airbyte@stroopwafelshop.iam.gserviceaccount.com	Enabled	airbyte	Used by Airbyte to connect to BigQuery	No keys		108688681614660056272	<div> <div></div> <div></div> <div></div> </div>

Actions

Figure 7 - The newly created service account

Creating a Service Account Key Pair

In Google Cloud, a service account key works much like a password, but it's actually a pair of cryptographic keys designed to authenticate your account with Google Cloud services. Essentially, when you want to connect to Google Cloud from any third-party application, you'll need to create this key file. Think of it as creating a special access pass that these applications can use to safely interact with Google Cloud on your behalf. It's crucial to handle this key file very carefully; if it falls into the wrong hands, they could gain unauthorized access to your Google Cloud resources. You can read more on service accounts keys here: <https://cloud.google.com/iam/docs/best-practices-for-managing-service-account-keys>

Now, navigate to the key creation page (as shown in **Figure 8**) and create a new **JSON** key type. Once it's created, you should download it to your local machine. Treat this with care! As a security measure, you cannot download or see the private key again. In case of a lost private key, you will have to create a new key pair and delete the old key.

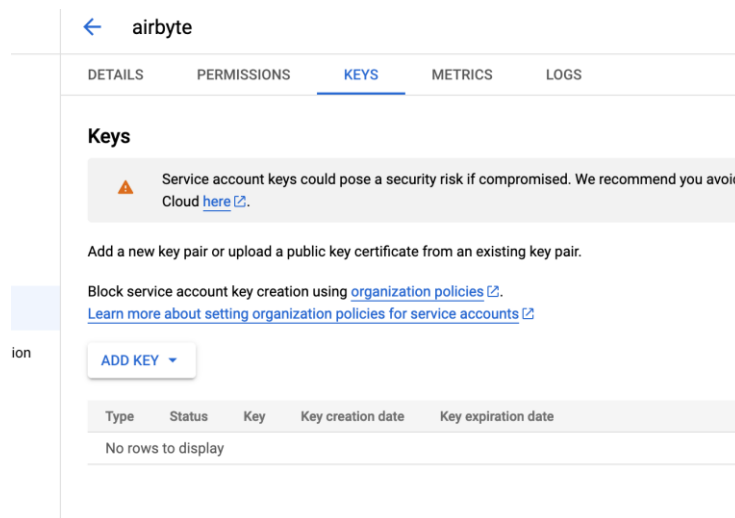


Figure 8 - Creating a service account key

You can open this file with a text-editor and take a look at the contents if you want. You will see that it contains a key-value mapping of different attributes such as **client_email** which corresponds to the service account's email, and most importantly the **private key**, which maps to the long string of letters and words. As you can see, the string does not make much sense to us humans. Nevertheless, it is important to keep the key secure and only hand it out to trusted parties. Next step: creating additional service accounts.

Creating additional service accounts

You just created a service account for Airbyte. During the Stroopwafelshop use case, we will also make use of **dbt Cloud** for modeling the data and **Tableau** for visualizing the results. Since it is a best practice to create separate service accounts for each third-party service, you should create two more service accounts: one for dbt Cloud, and one for Tableau.

For the dbt service account, choose **dbt-cloud** as the name, and assign it the **BigQuery Data Editor** and **BigQuery User** roles during the creation process. Then generate a key and download it, like you did previously. Next, for the Tableau service account, choose **tableau** as the name, and assign it the **BigQuery Data Viewer** role for read-only access. Finally, create the key file and download it. Return to the **IAM** overview. It should list your email address and the three service accounts with their assigned roles, as shown in **Figure 13**.

The screenshot shows the Google Cloud IAM & Admin console. The left sidebar contains the navigation menu with 'IAM & Admin' selected. The main content area is titled 'IAM' and shows 'Permissions for project "stroopwafelshop"'. Below this, there are tabs for 'PERMISSIONS' and 'RECOMMENDATIONS HISTORY'. A message at the top states: 'Beginning on April 29th, 2024 at-scale policy analysis and advanced IAM recommendation capabilities will require Security Command Center Premium. [Learn more](#)'. The 'VIEW BY PRINCIPALS' tab is selected, and the 'GRANT ACCESS' button is visible. A table lists the principals and their roles:

Type	Principal	Name	Role	Security Insights
Person	[redacted].com	ae_book_tester	Owner	9149/9211 excess permissions
Service Account	airbyte@stroopwafelshop.iam.gserviceaccount.com	airbyte	BigQuery User	25/25 excess permissions
Service Account	dbt-cloud@stroopwafelshop.iam.gserviceaccount.com	dbt-cloud	BigQuery Data Owner	52/65 excess permissions
Service Account			BigQuery Job User	2/4 excess permissions
Service Account	tableau@stroopwafelshop.iam.gserviceaccount.com	tableau	BigQuery Data Viewer	

Figure 9 – Three service accounts

Equipped with your service accounts, you are now ready to load data into BigQuery and interact with the third-party services. Continue where you left off in **Chapter 8 - Hands-on Analytics Engineering**.

